

C2_W4_Assignment

May 26, 2025

IMPORTANT

Please use the **Submit Assignment** button on the top (marked as 1) to submit your assignment for grading. **Do not** use the **Submit Assignment** button displayed on the right (marked as 2).

If you do not see the **Submit Assignment** button marked as 1, save your assignment and then re-open it.

1 Cox Proportional Hazards and Random Survival Forests

Welcome to the final assignment in Course 2! In this assignment you'll develop risk models using survival data and a combination of linear and non-linear techniques. We'll be using a dataset with survival data of patients with Primary Biliary Cirrhosis (pbc). PBC is a progressive disease of the liver caused by a buildup of bile within the liver (cholestasis) that results in damage to the small bile ducts that drain bile from the liver. Our goal will be to understand the effects of different factors on the survival times of the patients. Along the way you'll learn about the following topics:

- Cox Proportional Hazards
 - Data Preprocessing for Cox Models.
- Random Survival Forests
 - Permutation Methods for Interpretation.

1.1 Table of Contents

- Section ??
- Section ??
- Section ??
- Section ??
 - Section ??
- Section ??
- Section ??
 - Section ??
- Section ??
 - Section ??
- Section ??
- Section ??

1. Import Packages

We'll first import all the packages that we need for this assignment.

- `sklearn` is one of the most popular machine learning libraries.
- `numpy` is the fundamental package for scientific computing in python.
- `pandas` is what we'll use to manipulate our data.
- `matplotlib` is a plotting library.
- `lifelines` is an open-source survival analysis library.

```
[1]: import sklearn
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from lifelines import CoxPHFitter
from lifelines.utils import concordance_index as cindex
from sklearn.model_selection import train_test_split

from util import load_data

from public_tests import *
```

2. Load the Dataset

Run the next cell to load the data.

```
[2]: df = load_data('data/pbc.csv')
```

3. Explore the Dataset

In the lecture videos `time` was in months, however in this assignment, `time` will be converted into years. Also notice that we have assigned a numeric value to `sex`, where `female = 0` and `male = 1`.

Next, familiarize yourself with the data and the shape of it.

```
[3]: print(df.shape)

# df.head() only outputs the top few rows
df.head()
```

(258, 19)

```
[3]:
```

	time	status	trt	age	sex	ascites	hepato	spiders	edema	\
0	1.095890	1.0	0.0	58.765229	0.0	1.0	1.0	1.0	1.0	
1	12.328767	0.0	0.0	56.446270	0.0	0.0	1.0	1.0	0.0	
2	2.772603	1.0	0.0	70.072553	1.0	0.0	0.0	0.0	0.5	
3	5.273973	1.0	0.0	54.740589	0.0	0.0	1.0	1.0	0.5	
6	5.019178	0.0	1.0	55.534565	0.0	0.0	1.0	0.0	0.0	

	bili	chol	albumin	copper	alk.phos	ast	trig	platelet	protime	\
--	------	------	---------	--------	----------	-----	------	----------	---------	---

0	14.5	261.0	2.60	156.0	1718.0	137.95	172.0	190.0	12.2
1	1.1	302.0	4.14	54.0	7394.8	113.52	88.0	221.0	10.6
2	1.4	176.0	3.48	210.0	516.0	96.10	55.0	151.0	12.0
3	1.8	244.0	2.54	64.0	6121.8	60.63	92.0	183.0	10.3
6	1.0	322.0	4.09	52.0	824.0	60.45	213.0	204.0	9.7

	stage
0	4.0
1	3.0
2	4.0
3	4.0
6	3.0

Take a minute to examine particular cases.

```
[4]: i = 20
      df.iloc[i, :]
```

```
[4]: time          11.175342
      status        1.000000
      trt           0.000000
      age          44.520192
      sex           1.000000
      ascites       0.000000
      hepato        1.000000
      spiders       0.000000
      edema         0.000000
      bili          2.100000
      chol          456.000000
      albumin       4.000000
      copper        124.000000
      alk.phos     5719.000000
      ast           221.880000
      trig          230.000000
      platelet      70.000000
      protime       9.900000
      stage         2.000000
      Name: 23, dtype: float64
```

Now, split your dataset into train, validation and test set using 60/20/20 split.

```
[5]: np.random.seed(0)
      df_dev, df_test = train_test_split(df, test_size = 0.2)
      df_train, df_val = train_test_split(df_dev, test_size = 0.25)

      print("Total number of patients:", df.shape[0])
      print("Total number of patients in training set:", df_train.shape[0])
      print("Total number of patients in validation set:", df_val.shape[0])
```

```
print("Total number of patients in test set:", df_test.shape[0])
```

Total number of patients: 258
 Total number of patients in training set: 154
 Total number of patients in validation set: 52
 Total number of patients in test set: 52

Before proceeding to modeling, let's normalize the continuous covariates to make sure they're on the same scale. Again, we should normalize the test data using statistics from the train data.

```
[6]: continuous_columns = ['age', 'bili', 'chol', 'albumin', 'copper', 'alk.phos',
    ↪ 'ast', 'trig', 'platelet', 'protime']
mean = df_train.loc[:, continuous_columns].mean()
std = df_train.loc[:, continuous_columns].std()
df_train.loc[:, continuous_columns] = (df_train.loc[:, continuous_columns] -
    ↪ mean) / std
df_val.loc[:, continuous_columns] = (df_val.loc[:, continuous_columns] - mean) /
    ↪ std
df_test.loc[:, continuous_columns] = (df_test.loc[:, continuous_columns] -
    ↪ mean) / std
```

Let's check the summary statistics on our training dataset to make sure it's standardized.

```
[7]: df_train.loc[:, continuous_columns].describe()
```

```
[7]:
```

	age	bili	chol	albumin	copper	\
count	1.540000e+02	1.540000e+02	1.540000e+02	1.540000e+02	1.540000e+02	
mean	9.833404e-16	-3.258577e-16	1.153478e-16	1.153478e-16	5.767392e-18	
std	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	
min	-2.304107e+00	-5.735172e-01	-1.115330e+00	-3.738104e+00	-9.856552e-01	
25%	-6.535035e-01	-4.895812e-01	-5.186963e-01	-5.697976e-01	-6.470611e-01	
50%	-6.443852e-03	-3.846612e-01	-2.576693e-01	5.663556e-02	-3.140636e-01	
75%	5.724289e-01	2.977275e-02	1.798617e-01	6.890921e-01	3.435366e-01	
max	2.654276e+00	5.239050e+00	6.243146e+00	2.140730e+00	5.495204e+00	

	alk.phos	ast	trig	platelet	protime
count	1.540000e+02	1.540000e+02	1.540000e+02	1.540000e+02	1.540000e+02
mean	1.326500e-16	-1.263059e-15	8.074349e-17	2.018587e-17	1.291896e-14
std	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00	1.000000e+00
min	-7.882167e-01	-1.489281e+00	-1.226674e+00	-2.058899e+00	-1.735556e+00
25%	-5.186471e-01	-8.353982e-01	-6.884514e-01	-6.399831e-01	-7.382590e-01
50%	-3.416086e-01	-2.260984e-01	-2.495932e-01	-4.100373e-02	-1.398807e-01
75%	-4.620597e-03	6.061159e-01	3.755727e-01	6.617988e-01	3.587680e-01
max	4.869263e+00	3.058176e+00	5.165751e+00	3.190823e+00	4.447687e+00

4. Cox Proportional Hazards

Our goal is to build a risk score using the survival data that we have. We'll begin by fitting a Cox Proportional Hazards model to your data.

Recall that the Cox Proportional Hazards model describes the hazard for an individual i at time t as

$$\lambda(t, x) = \lambda_0(t)e^{\theta^T X_i}$$

The λ_0 term is a baseline hazard and incorporates the risk over time, and the other term incorporates the risk due to the individual's covariates. After fitting the model, we can rank individuals using the person-dependent risk term $e^{\theta^T X_i}$.

Categorical variables cannot be used in a regression model as they are. In order to use them, conversion to a series of variables is required.

Since our data has a mix of categorical (**stage**) and continuous (**wb1c**) variables, before we proceed further we need to do some data engineering. To tackle the issue at hand we'll be using the **Dummy Coding** technique. In order to use Cox Proportional Hazards, we will have to turn the categorical data into one hot features so that we can fit our Cox model. Luckily, Pandas has a built-in function called **get_dummies** that will make it easier for us to implement our function. It turns categorical features into multiple binary features.

Exercise 1 - to_one_hot In the cell below, implement the **to_one_hot(...)** function.

Hints

Remember to drop the first dummy for each each category to avoid convergence issues when fitting the proportional hazards model.

Check out the **get_dummies()** documentation.

Use **dtype=np.float64**.

```
[8]: # UNQ_C1 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def to_one_hot(dataframe, columns):
    """
    Convert columns in dataframe to one-hot encoding.
    Args:
        dataframe (dataframe): pandas dataframe containing covariates
        columns (list of strings): list categorical column names to one hot_
    → encode
    Returns:
        one_hot_df (dataframe): dataframe with categorical columns encoded
        as binary variables
    """

    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###

    one_hot_df = pd.get_dummies(dataframe, columns=columns, drop_first = True,
    → dtype=np.float64)

    ### END CODE HERE ###

    return one_hot_df
```

```
[9]: # test cell ex1 - do not modify this test cell
to_one_hot_test(to_one_hot, df_train, df_val, df_test)
```

One hot val columns:

```
['time', 'status', 'trt', 'age', 'sex', 'ascites', 'hepato', 'edema', 'bili',
'chol', 'albumin', 'copper', 'alk.phos', 'ast', 'trig', 'platelet', 'protime',
'spiders_1.0', 'stage_2.0', 'stage_3.0', 'stage_4.0']
```

There are 21 columns

All tests passed.

Expected Output:

One hot val columns:

```
['time', 'status', 'trt', 'age', 'sex', 'ascites', 'hepato', 'edema', 'bili', 'chol', 'albumin
```

There are 21 columns

All tests passed.

Now you'll use the function you coded to transform the training, validation, and test sets.

```
[10]: to_encode = ['edema', 'stage']

one_hot_train = to_one_hot(df_train, to_encode)
one_hot_val = to_one_hot(df_val, to_encode)
one_hot_test = to_one_hot(df_test, to_encode)
```

1.1.1 Look for new features

Now, let's take a peek at one of the transformed data sets. Do you notice any new features?

```
[11]: print(one_hot_train.shape)
one_hot_train.head()
```

(154, 22)

```
[11]:
```

	time	status	trt	age	sex	ascites	hepato	spiders	\
279	3.868493	0.0	0.0	-0.414654	0.0	0.0	0.0	0.0	
137	3.553425	1.0	0.0	0.069681	1.0	0.0	1.0	0.0	
249	4.846575	0.0	1.0	-0.924494	0.0	0.0	1.0	0.0	
266	0.490411	1.0	0.0	1.938314	0.0	1.0	1.0	1.0	
1	12.328767	0.0	0.0	0.563645	0.0	0.0	1.0	1.0	

	bili	chol	...	alk.phos	ast	trig	platelet	\
279	-0.300725	-0.096081	...	0.167937	0.401418	0.330031	0.219885	
137	0.895363	0.406085	...	0.101665	0.472367	1.621764	-0.120868	
249	-0.510565	-0.225352	...	0.245463	1.899020	-0.580807	0.422207	
266	0.748475	-0.608191	...	-0.650254	-0.288898	-0.481443	-0.727833	
1	-0.405645	-0.210436	...	2.173526	-0.144699	-0.531125	-0.450972	

	protime	edema_0.5	edema_1.0	stage_2.0	stage_3.0	stage_4.0
279	-1.137178	0.0	0.0	0.0	1.0	0.0
137	-0.239610	0.0	0.0	0.0	1.0	0.0
249	0.159309	0.0	0.0	0.0	0.0	1.0
266	1.356065	0.0	1.0	0.0	0.0	1.0
1	-0.139881	0.0	0.0	0.0	1.0	0.0

[5 rows x 22 columns]

5. Fitting and Interpreting a Cox Model

Run the following cell to fit your Cox Proportional Hazards model using the `lifelines` package.

```
[12]: cph = CoxPHFitter()
      cph.fit(one_hot_train, duration_col = 'time', event_col = 'status', step_size=0.
      ↪1)
```

```
[12]: <lifelines.CoxPHFitter: fitted with 154 total observations, 90 right-censored
      observations>
```

You can use `cph.print_summary()` to view the coefficients associated with each covariate as well as confidence intervals.

```
[13]: cph.print_summary()
```

```
<IPython.core.display.HTML object>
```

Question:

- According to the model, was treatment `trt` beneficial?
- What was its associated hazard ratio?
 - Note that the hazard ratio is how much an incremental increase in the feature variable changes the hazard.


Check your answer!

You should see that the treatment (`trt`) was beneficial because it has a negative impact on the hazard (the coefficient is negative, and `exp(coef)` is less than 1).

The associated hazard ratio is ~0.8, because this is the `exp(coef)` of treatment.

We can compare the predicted survival curves for treatment variables. Run the next cell to plot survival curves using the `plot_covariate_groups()` function. - The y-axis is the survival rate - The x-axis is time

```
[14]: cph.plot_covariate_groups('trt', values=[0, 1]);
```



Notice how the group without treatment has a lower survival rate at all times (the x-axis is time) compared to the treatment group.

6. Hazard Ratio

Recall from the lecture videos that the Hazard Ratio between two patients was the likelihood of one patient (e.g smoker) being more at risk than the other (e.g non-smoker).

$$\frac{\lambda_{smoker}(t)}{\lambda_{nonsmoker}(t)} = e^{\theta(X_{smoker} - X_{nonsmoker})^T}$$

Where

$$\lambda_{smoker}(t) = \lambda_0(t)e^{\theta X_{smoker}^T}$$

and

$$\lambda_{nonsmoker}(t) = \lambda_0(t)e^{\theta X_{nonsmoker}^T}$$

Exercise 2 - hazard_ratio In the cell below, write a function to compute the hazard ratio between two individuals given the cox model's coefficients.

Hints

use numpy.dot

use numpy.exp


```
[17]: # UNQ_C2 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def hazard_ratio(case_1, case_2, cox_params):
    '''
    Return the hazard ratio of case_1 : case_2 using
    the coefficients of the cox model.

    Args:
        case_1 (np.array): (1 x d) array of covariates
        case_2 (np.array): (1 x d) array of covariates
        model (np.array): (1 x d) array of cox model coefficients
    Returns:
        hazard_ratio (float): hazard ratio of case_1 : case_2
    '''

    ### START CODE HERE (REPLACE INSTANCES OF 'None' with your code) ###

    hr = np.exp(np.dot(cox_params, (case_1 - case_2).T))

    ### END CODE HERE ###

    return hr
```

Now, evaluate it on the following pair of individuals: $i = 1$ and $j = 5$. Given your implementation of `hazard_ratio` is correct, you will **only** pass the test when $i = 1$ and $j = 5$.

```
[18]: # test cell ex2

# Set i = 1 and j = 5 to pass the test
i = 1
j = 5

# DO NOT modify anything below
hazard_ratio_test(hazard_ratio, i, j, one_hot_train, cph)
```

15.029017732492221

All tests passed.

Expected Output:

15.029017732492221

All tests passed.

Question:

Is `case_1` or `case_2` at greater risk?

Check your answer!

Important! The following answer only applies if you picked $i = 1$ and $j = 5$

You should see that `case_1` is at higher risk.

The hazard ratio of case 1 / case 2 is greater than 1, so case 1 had a higher hazard relative to case 2

Inspect different pairs, and see if you can figure out which patient is more at risk.

```
[19]: i = 4
      case_1 = one_hot_train.iloc[i, :].drop(['time', 'status'])

      j = 7
      case_2 = one_hot_train.iloc[j, :].drop(['time', 'status'])

      print("Case 1\n\n", case_1, "\n")
      print("Case 2\n\n", case_2, "\n")
      print("Hazard Ratio:", hazard_ratio(case_1.values, case_2.values, cph.params_
      ↪values))
```

Case 1

trt	0.000000
age	0.563645
sex	0.000000
ascites	0.000000
hepato	1.000000
spiders	1.000000
bili	-0.405645
chol	-0.210436
albumin	1.514297
copper	-0.481961
alk.phos	2.173526
ast	-0.144699
trig	-0.531125
platelet	-0.450972
protime	-0.139881
edema_0.5	0.000000
edema_1.0	0.000000
stage_2.0	0.000000
stage_3.0	1.000000
stage_4.0	0.000000

Name: 1, dtype: float64

Case 2

trt	0.000000
age	0.463447
sex	0.000000

```

ascites      0.000000
hepato       1.000000
spiders      0.000000
bili        -0.489581
chol        -0.309875
albumin     -1.232371
copper      -0.504348
alk.phos    2.870427
ast         -0.936261
trig        -0.150229
platelet     3.190823
protime     -0.139881
edema_0.5    0.000000
edema_1.0    0.000000
stage_2.0    0.000000
stage_3.0    0.000000
stage_4.0    1.000000
Name: 38, dtype: float64

```

Hazard Ratio: 0.1780450006997129

Check your answer!

Important! The following answer only applies if you picked $i = 4$ and $j = 7$

You should see that `case_2` is at higher risk.

The hazard ratio of case 1 / case 2 is less than 1, so case 2 had a higher hazard relative to case 1

7. Harrell's C-index

To evaluate how good our model is performing, we will write our own version of the C-index. Similar to the week 1 case, C-index in the survival context is the probability that, given a randomly selected pair of individuals, the one who died sooner has a higher risk score.

However, we need to take into account censoring. Imagine a pair of patients, A and B .

Scenario 1

- A was censored at time t_A
- B died at t_B
- $t_A < t_B$.

Because of censoring, we can't say whether A or B should have a higher risk score.

Scenario 2 Now imagine that $t_A > t_B$.

- A was censored at time t_A
- B died at t_B
- $t_A > t_B$

Now we can definitively say that B should have a higher risk score than A , since we know for a fact that A lived longer.

Therefore, when we compute our C-index - We should only consider pairs where at most one person is censored - If they are censored, then their censored time should occur *after* the other person's time of death.

The metric we get if we use this rule is called **Harrel's C-index**.

Note that in this case, being censored at time t means that the true death time was some time AFTER time t and not at t . - Therefore if $t_A = t_B$ and A was censored: - Then A actually lived longer than B . - This will effect how you deal with ties in the exercise below!

Exercise 3 - harrell_c Fill in the function below to compute Harrel's C-index.

Hints

If you get a division by zero error, consider checking how you count when a pair is permissible (in the case where one patient is censored and the other is not censored).

```
[20]: # UNQ_C3 (UNIQUE CELL IDENTIFIER, DO NOT EDIT)
def harrell_c(y_true, scores, event):
    """
    Compute Harrel C-index given true event/censoring times,
    model output, and event indicators.

    Args:
        y_true (array): array of true event times
        scores (array): model risk scores
        event (array): indicator, 1 if event occurred at that index, 0 for
        →censorship
    Returns:
        result (float): C-index metric
    """

    n = len(y_true)
    assert (len(scores) == n and len(event) == n)

    concordant = 0.0
    permissible = 0.0
    ties = 0.0

    result = 0.0

    ### START CODE HERE (REPLACE INSTANCES OF 'None' and 'pass' with your code)
    →###

    # use double for loop to go through cases
    for i in range(n):
        # set lower bound on j to avoid double counting
```

```

for j in range(i+1, n):

    # check if at most one is censored
    if not (event[i] == 0 and event[j] == 0):

        # check if neither are censored
        if event[i] == 1 and event[j] == 1:
            permissible += 1

        # check if scores are tied
        if y_true[i] == y_true[j]:
            ties += 1

        # check for concordant
        elif y_true[i] > y_true[j] and scores[i] < scores[j]:
            concordant += 1
        elif y_true[i] < y_true[j] and scores[i] > scores[j]:
            concordant += 1

    # check if one is censored
    elif event[i] == 0 or event[j] == 0:

        # get censored index
        censored = j
        uncensored = i

        if event[i] == 0:
            censored = i
            uncensored = j

        # check if permissible
        # Note: in this case, we are assuming that censored at a
        ↪ time

        # means that you did NOT die at that time. That is, if you
        # live until time 30 and have event = 0, then you lived
        ↪ THROUGH

        # time 30.
        if y_true[censored] >= y_true[uncensored]:
            permissible += 1

        # check if scores are tied
        if scores[i] == scores[j]:
            # update ties
            ties += 1

        # check if scores are concordant

```

```

        if y_true[censored] >= y_true[uncensored] and
→scores[censored] < scores[uncensored]:
            concordant += 1

    # set result to c-index computed from number of concordant pairs,
    # number of ties, and number of permissible pairs (REPLACE 0 with your
→code)
    result = (concordant + 0.5 * ties) / permissible

    ### END CODE HERE ###

    return result

```

```

[21]: # test cell ex1 - do not modify this test cell
      harrell_c_test(harrell_c)

```

Test Case 1

```

y_true:  [30, 12, 84, 9]
scores:  [0.5, 0.9, 0.1, 1.0]
event:    [1, 1, 1, 1]
Output:  1.0

```

Test Case 2

```

y_true:  [30, 12, 84, 9]
scores:  [0.9, 0.5, 1.0, 0.1]
event:    [1, 1, 1, 1]
Output:  0.0

```

Test Case 3

```

y_true:  [30, 12, 84, 9]
scores:  [0.5, 0.9, 0.1, 1.0]
event:    [1, 0, 1, 1]
Output:  1.0

```

Test Case 4

```

y_true:  [30, 30, 20, 20]
scores:  [10, 5, 15, 20]
event:    [1, 0, 1, 0]
Output:  0.75

```

Test Case 5

```

y_true:  [20, 20, 30, 30, 30]

```

```
scores: [20, 15, 5, 10, 15]
event:  [0, 1, 0, 1, 0]
Output: 0.5833333333333334
```

Tes Case 6

```
y_true: [10, 10]
scores: [4, 5]
event:  [0, 1]
Output: 1.0
```

All tests passed.

Expected Output:

Test Case 1
Output: 1.0

Test Case 2
Output: 0.0

Test Case 3
Output: 1.0

Test Case 4
Output: 0.75

Test Case 5
Output: 0.5833333333333334

Tes Case 6
Output: 1.0

All tests passed.

Now use the Harrell's C-index function to evaluate the cox model on our data sets.

```
[22]: # Train
scores = cph.predict_partial_hazard(one_hot_train)
cox_train_scores = harrell_c(one_hot_train['time'].values, scores.values,
    ↪one_hot_train['status'].values)
# Validation
scores = cph.predict_partial_hazard(one_hot_val)
cox_val_scores = harrell_c(one_hot_val['time'].values, scores.values,
    ↪one_hot_val['status'].values)
# Test
scores = cph.predict_partial_hazard(one_hot_test)
```

```

cox_test_scores = harrell_c(one_hot_test['time'].values, scores.values,
    ↪one_hot_test['status'].values)

print("Train:", cox_train_scores)
print("Val:", cox_val_scores)
print("Test:", cox_test_scores)

```

```

Train: 0.8265139116202946
Val: 0.8544776119402985
Test: 0.8478543563068921

```

What do these values tell us ?

8. Random Survival Forests

This performed well, but you have a hunch you can squeeze out better performance by using a machine learning approach. You decide to use a Random Survival Forest. To do this, you can use the `RandomForestSRC` package in R. To call R function from Python, we'll use the `r2py` package. Run the following cell to import the necessary requirements.

```

[23]: %load_ext rpy2.ipynon
      %R require(ggplot2)

      from rpy2.robjects.packages import importr
      # import R's "base" package
      base = importr('base')

      # import R's "utils" package
      utils = importr('utils')

      # import rpy2's package module
      import rpy2.robjects.packages as rpackages

      forest = rpackages.importr('randomForestSRC', lib_loc='R')

      from rpy2 import robjects as ro
      R = ro.r

      from rpy2.robjects import pandas2ri
      pandas2ri.activate()

```

```
R[write to console]: Loading required package: ggplot2
```

Instead of encoding our categories as binary features, we can use the original dataframe since trees deal well with raw categorical data (can you think why this might be?).

Run the code cell below to build your forest.


```
[24]: model = forest.rfsrc(ro.Formula('Surv(time, status) ~ .'), data=df_train,
    ↪ ntree=300, nodedepth=5, seed=-1)
```

```
[25]: print(model)
```

```

                Sample size: 154
            Number of deaths: 64
            Number of trees: 300
        Forest terminal node size: 15
    Average no. of terminal nodes: 6.54
No. of variables tried at each split: 5
        Total no. of variables: 17
    Resampling used to grow trees: swor
    Resample size used to grow trees: 97
            Analysis: RSF
            Family: surv
        Splitting rule: logrank *random*
    Number of random split points: 10
            Error rate: 19.07%
```

Finally, let's evaluate on our validation and test sets, and compare it with our Cox model.

```
[26]: result = R.predict(model, newdata=df_val)
    scores = np.array(result.rx('predicted')[0])

    print("Cox Model Validation Score:", cox_val_scores)
    print("Survival Forest Validation Score:", harrell_c(df_val['time'].values,
    ↪ scores, df_val['status'].values))
```

```
Cox Model Validation Score: 0.8544776119402985
Survival Forest Validation Score: 0.8296019900497512
```

```
[27]: result = R.predict(model, newdata=df_test)
    scores = np.array(result.rx('predicted')[0])

    print("Cox Model Test Score:", cox_test_scores)
    print("Survival Forest Validation Score:", harrell_c(df_test['time'].values,
    ↪ scores, df_test['status'].values))
```

```
Cox Model Test Score: 0.8478543563068921
Survival Forest Validation Score: 0.8621586475942783
```

Your random forest model should be outperforming the Cox model slightly. Let's dig deeper to see how they differ.

9. Permutation Method for Interpretation


We'll dig a bit deeper into interpretation methods for forests a bit later, but for now just know that

random survival forests come with their own built in variable importance feature. The method is referred to as VIMP, and for the purpose of this section you should just know that higher absolute value of the VIMP means that the variable generally has a larger effect on the model outcome.

Run the next cell to compute and plot VIMP for the random survival forest.

```
[28]: vimps = np.array(forest.vimp(model).rx('importance')[0])

y = np.arange(len(vimps))
plt.barh(y, np.abs(vimps))
plt.yticks(y, df_train.drop(['time', 'status'], axis=1).columns)
plt.title("VIMP (absolute value)")
plt.show()
```



1.1.2 Question:

How does the variable importance compare to that of the Cox model? Which variable is important in both models? Which variable is important in the random survival forest but not in the Cox model? You should see that **edema** is important in both the random survival forest and the Cox model. You should also see that **bili** is important in the random survival forest but not the Cox model.

1.2 Congratulations!

You've finished the last assignment in course 2! Take a minute to look back at the analysis you've done over the last four assignments. You've done a great job!